



Per-context predictions behind the five corrected sycophancy rungs (fair-allocation refit)


model-written evaluations (n=576 contexts, seed 0)




answer (n=594 contexts, seed 0)




feedback (n=559 contexts, seed 0)



are-you-sure (n=598 contexts, seed 0)



mimicry (n=285 contexts, seed 0)



judge-scored sycophancy DV (per-context mean)

judge-scored sycophancy DV (per-context mean)