

Map reconstruction and retrieval ON the behavior evaluation distributions

Map fit on 15,034 WildChat context->answer pairs (U = 18,793 refit), context end state. Left: reconstruction R^2 per layer, in-distribution vs on each evaluation rung. Middle / right: kNN retrieval of the true target answer state among that rung's own held-out pool, best layer, cosine metric; the black dash is the chance rate (k / n_{pool}) and the annotation is the multiple of chance. Retrieval stays far above chance where R^2 is strongly negative.

