

Out-of-distribution scatters: predicted score vs judged expression

One point per evaluation context (a random 1,500-point subsample is drawn for legibility; the annotated Spearman rho and n are computed on ALL points).
U = 18,793 unlabeled map pairs, E1, context end state, at the largest label budget the persisted predictions cover per rung (stated per panel).
DV construct: evil / sycophancy graded 0-100 trait score; hallucination 0-1 fabrication rate on its own rungs.

