

R5: replacing half the unlabeled map pool with trait-eliciting contexts

Held-out train rung, context end state, E1, each behavior at its MAXIMUM label budget (evil L=8,000; sycophancy / hallucination L=16,000).

Blue = all-generic WildChat map pool; orange = same pool size with 50% of the pairs replaced by UNLABELED trait-eliciting contexts.

Error bars are the 95% bootstrap CI over contexts from each row's ci_frozen. Bottom row is the built-in control: direct ridge never reads the unlabeled map pool, so any movement there is labeled-draw resampling, not a pool effect.

