





#1739 — persona-vector readouts by evaluation setting (P-A)

P-A readout: trained on one trait-eliciting dataset + judged WildChat split | MLP-map arm from the earlier nonlinear_map round (generic-pool map fit; no pvsynth/WildChat rungs)



■ PV on context
 ▨ PV on mapped answer — MLP map
 △ low DV spread — untrustworthy
■ PV on mapped answer — linear map
 ■ PV on real answer (upper bound)