

#1739 — persona-vector readouts by evaluation setting (P-A)

P-A readout: trained on one trait-eliciting dataset + judged WildChat split | MLP-map arm from the earlier nonlinear_map round (generic-pool map fit; no pvsynth/WildChat rungs)

