


Query -> answer map: does the bare-query readout improve with more labels?

sycophancy


n = 1,982 eval contexts (WildChat rung)

Spearman ρ vs judged DV (frozen layer)



hallucination

n = 1,967 eval contexts (WildChat rung)



evil

n = 1,987 eval contexts (WildChat rung)

