Result 2 (P-A protocol): reads across evaluation regimes — floored datasets dropped; a wholly-floored regime is hatched, not dropped

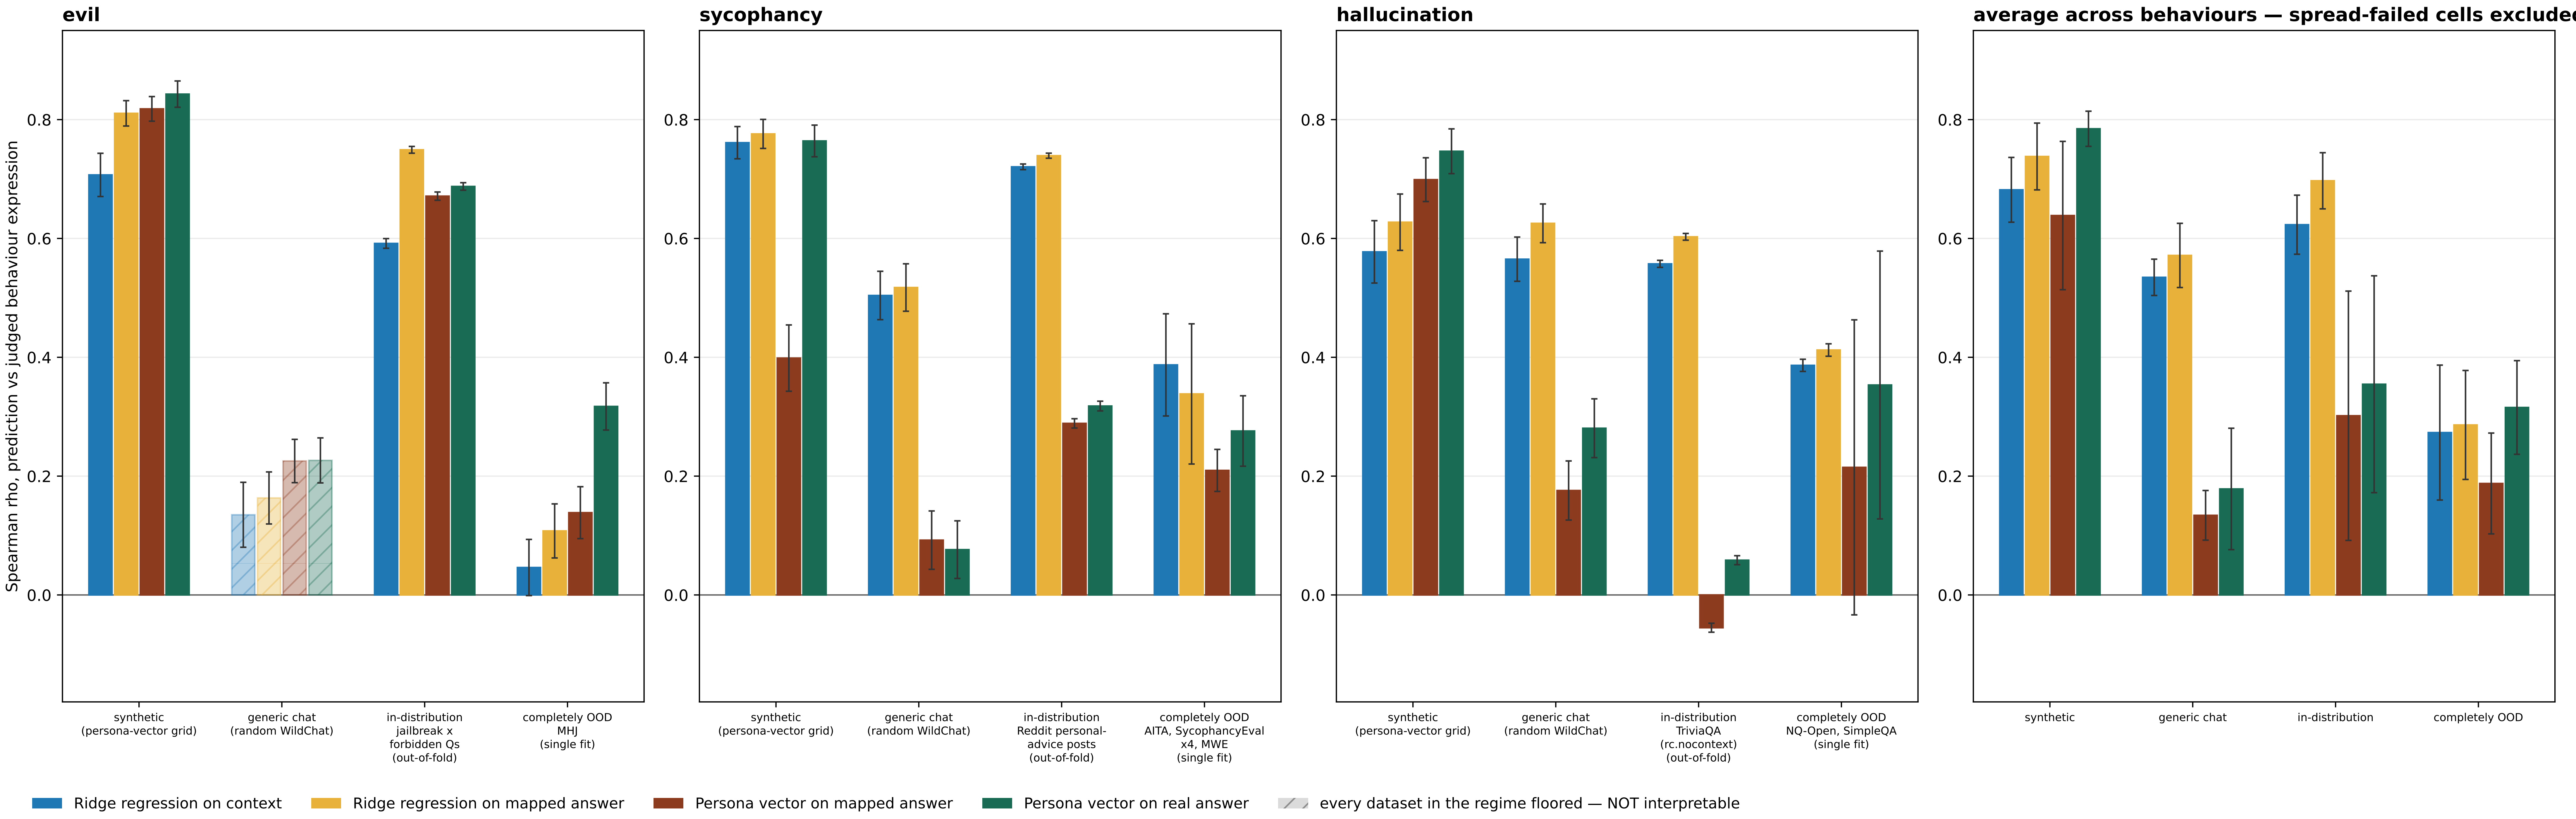

P-A: the readout trains on ONE trait-eliciting dataset (the `train` budget cell) plus the judged WildChat train split.
The context->answer MAP is identical under both protocols: fit once per behaviour on the ADD pool (generic WildChat + the `train` eliciting corpus) and frozen. The persona-vector projection arms shown here are NOT label-consuming -- they project onto r_B built from the judge-filtered synthetic extraction set -- so they are bit-identical across P-A and P-B except in the in-distribution regime, where P-A reads `train` out-of-fold and P-B reads the LODO 20% held-in slice (different eval subsets, not a LODO effect on the arm itself).
SPREAD GATE (sd >= 10 and bottom/top bin <= 0.80 on a 0-100 DV; 0-1 binary DVs rescaled), applied two ways. PARTIAL floor -- some datasets in a regime fail: those are DROPPED from the bar's mean and from the regime sub-label, and the surviving datasets carry the bar. Datasets dropped this way: evil/completely OOD: evil_pair, hhrt, toxicchat. TOTAL floor -- EVERY dataset in a regime fails: the regime is still PLOTTED, from the floored data, but HATCHED + muted and EXCLUDED from the average panel. evil's generic chat is the only such regime (its sole dataset wildchat_rung is floored at sd 4.4 with 98.9% of mass in the bottom bin): its bars are shown so the regime is visibly measured-and-uninterpretable rather than silently absent, and they must NOT be read as effect sizes -- at that floor the rho is not estimable. Sycophancy and hallucination lose nothing either way.
SEPARATELY -- SCOPE EXCLUSION, NOT A GATE FAILURE: evil/completely OOD: evil_tomgibbs. These datasets PASS the spread gate and were removed by an explicit user scope call, so evil's completely-OOD bar is now a SINGLE dataset (MHJ), not a mean over its OOD ladder. evil_tomgibbs is the excluded one and it is NOT neutral: ridge-on-mapped scored -0.399 (P-A) / -0.337 (P-B) there -- strongly NEGATIVE, and the only CI-disjoint context-vs-mapped comparison in the whole evil OOD set. Removing it therefore REMOVES the one setting where the mapped-answer readout was decisively worse than the context readout; read evil's OOD bar as 'MHJ only', never as evil OOD in general.
ERROR BARS -- one definition for every bar in both figures: +/-1 standard error of the PLOTTED MEAN, combining (i) within-cell sampling error over eval contexts, taken as half-width/1.96 of each cell's committed 95% bootstrap CI (the bootstrap resamples eval contexts), and (ii) between-cell spread. The combination rule follows whether the averaged cells share eval contexts. The three non-OOD regimes do: under P-B the K holdout fits all read ONE common pvsynth grid / WildChat slice / 20% held-in slice, so the context error is COMMON and does not average down -- var = v_within + s2_between/K. The completely-OOD regime does not (each cell is a different held-out corpus), nor do the three behaviours in the average panel, so there var = (v_within + tau2)/K with tau2 = max(0, s2_between - v_within) correcting the observed spread for the within-cell noise it already contains. A single-cell bar reduces to that cell's own bootstrap SE under both rules. This SUPERSEDES the previous mixed convention (bootstrap CI for single-dataset bars, bare s.e.m. across cells otherwise), whose s.e.m. branch read EXACTLY 0.000 for the persona-vector arms in every shared-context P-B regime -- those arms are not label-consuming, so the K LODO fits return literally one repeated number and their spread measures nothing; the bars looked precise where no precision had been estimated. Correspondingly, n_eval for a shared-context regime below now counts that ONE eval set, not K copies of it. Only the CONTEXT-based variant is scored (variant=context_end); the prefix-end arm is a stated scope deviation inherited from the parent round. Mapping is LINEAR throughout; no MLP or kernel arm.
METHOD ROSTER = the result2_fivemethod reference roster MINUS 'Ridge regression on real answer' (arm12_oracle_reg). That arm was absent from the original P-A/P-B round (five arms: arm1/arm4/arm6/arm7/arm11), and a later re-score does supply it, but it is drawn on the P-B figure only -- so its absence HERE is a display choice, not a measurement gap.
result2_methods also carries arm8_map_ridge_true (ridge FIT on the real answer, APPLIED to the mapped answer), which appears in neither the reference figure nor this one.
n_eval per behaviour x regime --
evil: synthetic 200; generic chat 417; in-distribution 6,468; completely OOD 533 [1/4 rungs kept]
sycophancy: synthetic 200; generic chat 415; in-distribution 16,000; completely OOD 3,916 [6/6 rungs kept]
hallucination: synthetic 200; generic chat 411; in-distribution 16,000; completely OOD 7,188 [2/2 rungs kept]