


Result 2 (P-A protocol), persona-vector spread-flagged view: reads across evaluation regimes, hatched where the DV spread gate fails



P-A: the readout trains on ONE trait-eliciting dataset (the `train` budget cell) plus the judged WildChat train split.  
The context->answer MAP is identical under both protocols: fit once per behaviour on the ADD pool (generic WildChat + the `train` eliciting corpus) and frozen. The persona-vector projection arms shown here are NOT label-consuming -- they project onto  $r_B$  built from the judge-filtered synthetic extraction set -- so they are bit-identical across P-A and P-B except in the in-distribution regime, where P-A reads `train` out-of-fold and P-B reads the 20% held-in slice (different eval subsets, not a LODO effect).  
Bars are the mean over spread-PASSING runs in the regime; error bars are the committed bootstrap CI for single-rung regimes and the s.e.m. across rungs for multi-rung ones. Hatched + muted = every rung in that cell FAILED the DV spread gate (sd  $\geq 1.0$  and bottom/top bin  $\leq 0.80$  on a 0-100 DV; 0-1 binary DVs rescaled) -- kept for completeness, not interpretable. Spread failures: evil hhrt / toxicchat / evil\_pair / wildchat\_rung.  
Only the CONTEXT-based variant is scored (variant=context\_end); the prefix-end arm is a stated scope deviation inherited from the parent round. Mapping is LINEAR throughout; no MLP or kernel arm.  
n\_eval per behaviour x regime --  
evil: synthetic 200; generic chat 417; in-distribution 6,468; completely OOD 2,033 [2/5 rungs kept]  
sycophancy: synthetic 200; generic chat 415; in-distribution 16,000; completely OOD 3,916 [6/6 rungs kept]  
hallucination: synthetic 200; generic chat 411; in-distribution 16,000; completely OOD 7,188 [2/2 rungs kept]