Result 2 (P-B protocol): reads across evaluation regimes — floored datasets dropped; a wholly-floored regime is hatched, not dropped

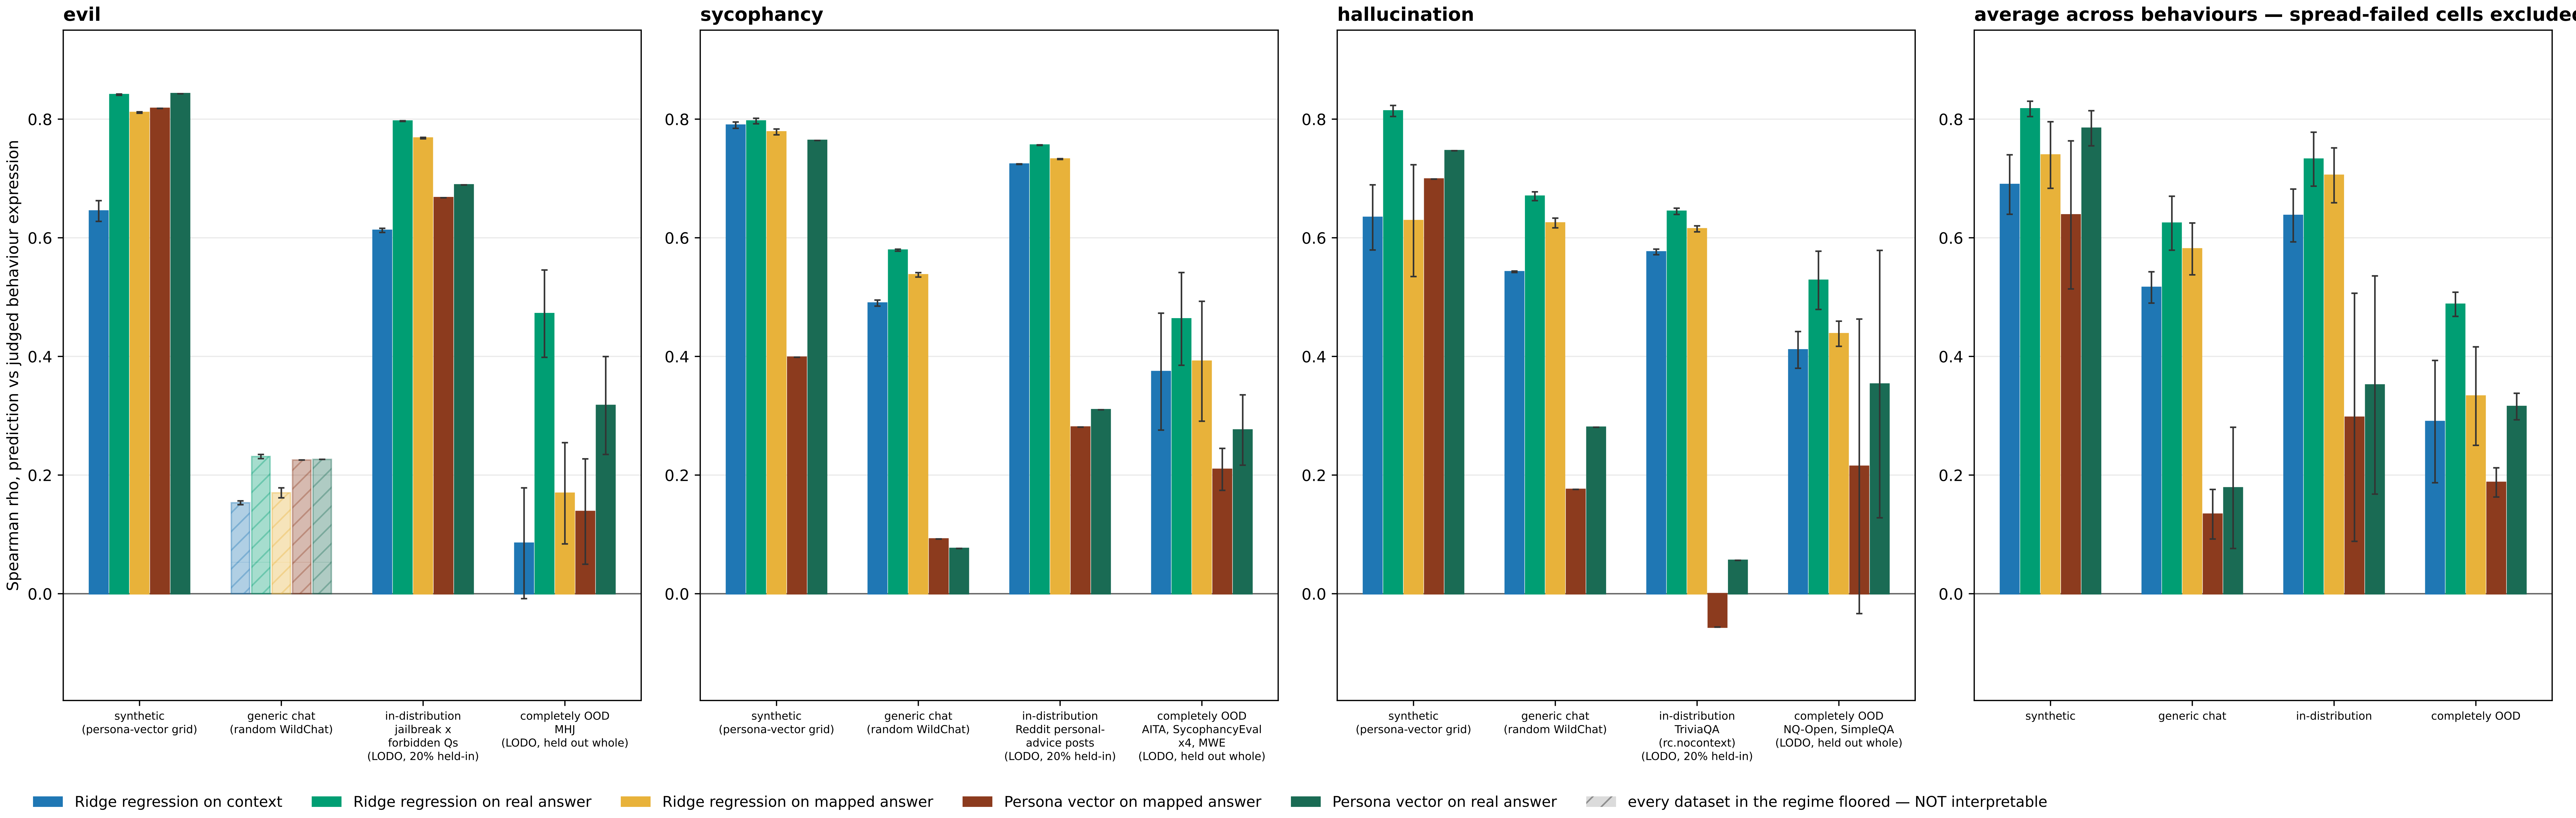

P-B (LODO): the readout trains on an 80% group-level slice of every trait-eliciting dataset EXCEPT one held out whole, plus the judged WildChat train split; one fit per holdout, and the 'completely OOD' bar is each fit's own held-out dataset.
The context->answer MAP is identical under both protocols: fit once per behaviour on the ADD pool (generic WildChat + the `train` eliciting corpus) and frozen. The persona-vector projection arms shown here are NOT label-consuming -- they project onto r_B built from the judge-filtered synthetic extraction set -- so they are bit-identical across P-A and P-B except in the in-distribution regime, where P-A reads `train` out-of-fold and P-B reads the LODO 20% held-in slice (different eval subsets, not a LODO effect on the arm itself).
SPREAD GATE (sd >= 10 and bottom/top bin <= 0.80 on a 0-100 DV; 0-1 binary DVs rescaled), applied two ways. PARTIAL floor -- some datasets in a regime fail: those are DROPPED from the bar's mean and from the regime sub-label, and the surviving datasets carry the bar. Datasets dropped this way: evil/completely OOD: hhrt, toxicchat, evil_pair. TOTAL floor -- EVERY dataset in a regime fails: the regime is still PLOTTED, from the floored data, but HATCHED + muted and EXCLUDED from the average panel. evil's generic chat is the only such regime (its sole dataset wildchat_rung is floored at sd 4.4 with 98.9% of mass in the bottom bin): its bars are shown so the regime is visibly measured-and-uninterpretable rather than silently absent, and they must NOT be read as effect sizes -- at that floor the rho is not estimable. Sycophancy and hallucination lose nothing either way.
SEPARATELY -- SCOPE EXCLUSION, NOT A GATE FAILURE: evil/completely OOD: evil_tomgibbs. These datasets PASS the spread gate and were removed by an explicit user scope call, so evil's completely-OOD bar is now a SINGLE dataset (MHJ), not a mean over its OOD ladder. evil_tomgibbs is the excluded one and it is NOT neutral: ridge-on-mapped scored -0.399 (P-A) / -0.337 (P-B) there -- strongly NEGATIVE, and the only CI-disjoint context-vs-mapped comparison in the whole evil OOD set. Removing it therefore REMOVES the one setting where the mapped-answer readout was decisively worse than the context readout; read evil's OOD bar as 'MHJ only', never as evil OOD in general.
!! ERROR BARS ARE THE SUPERSEDED (WRONG) DEFINITION -- this figure exists ONLY to show what the bars looked like before cad915214a. DO NOT use it for inference or publication; the calibrated twin is the same filename without the _oldci suffix. The rule restored here picks one of two branches BY CELL COUNT, and the two are not in the same units: a single-cell bar gets that cell's FULL 95% bootstrap half-width (~1.96 SE), while a multi-cell bar gets a bare s.e.m. across the cell MEANS, DISCARDING each cell's own bootstrap CI entirely -- so the multi-cell branch measures only how much the cells disagree with each other, treating every cell mean as exact. Measured on THIS figure's own 80 plotted bars, against the calibrated definition: the 5 single-cell bars are 2.03x TOO WIDE; of the 75 multi-cell bars, 18 read EXACTLY 0.000 (an invisible bar) and 55 are too NARROW, by a median factor of 1.35x and up to 30.4x. The exact zeros are the shared-context regimes, where the K LODO fits read ONE common eval set and the non-label-consuming arms return literally one repeated number, so their spread is 0 by construction. The understatement is systematic rather than occasional because the between-cell spread is almost always SMALLER than the within-cell bootstrap noise this rule omits. n_eval below is still the CORRECTED shared-context count (ONE eval set, not K copies); only the error rule was reverted.
Only the CONTEXT-based variant is scored (variant=context_end); the prefix-end arm is a stated scope deviation inherited from the parent round. Mapping is LINEAR throughout; no MLP or kernel arm.
METHOD ROSTER = the result2_fivemethod reference roster, all five arms. 'Ridge regression on real answer' (arm12_oracle_reg) was NOT part of the original P-A/P-B round -- that round scored five arms (arm1/arm4/arm6/arm7/arm11) and nothing ridge-on-the-real-answer -- so its bars here come from a SEPARATE re-score that re-ran the SAME P-A/P-B fits with the roster extended by that one arm (behaviours re-scored: evil, hallucination, sycophancy). The other four arms still read the committed round and are numerically UNCHANGED by the re-score, so no bar that was already published has moved; a reproduction gate checks that the five arms the two sources share came back identical, and that is what licenses drawing them side by side.
result2_methods also carries arm8_map_ridge_true (ridge FIT on the real answer, APPLIED to the mapped answer), which appears in neither the reference figure nor this one.
n_eval per behaviour x regime --
evil: synthetic 200 [5/5 rungs kept]; generic chat 417 [5/5 rungs kept]; in-distribution 1,162 [5/5 rungs kept]; completely OOD 533 [1/4 rungs kept]
sycophancy: synthetic 200 [6/6 rungs kept]; generic chat 415 [6/6 rungs kept]; in-distribution 3,235 [6/6 rungs kept]; completely OOD 3,916 [6/6 rungs kept]
hallucination: synthetic 200 [2/2 rungs kept]; generic chat 411 [2/2 rungs kept]; in-distribution 3,350 [2/2 rungs kept]; completely OOD 7,188 [2/2 rungs kept]