


Result 2 (P-B protocol): reads across evaluation regimes — floored datasets dropped; a wholly-floored regime is hatched, not dropped



P-B (LODO): the readout trains on an 80% group-level slice of every trait-eliciting dataset EXCEPT one held out whole, plus the judged WildChat train split; one fit per holdout, and the 'completely OOD' bar is each fit's own held-out dataset.

The context->answer MAP is identical under both protocols: fit once per behaviour on the ADD pool (generic WildChat + the `train` eliciting corpus) and frozen. The persona-vector projection arms shown here are NOT label-consuming -- they project onto r_B built from the judge-filtered synthetic extraction set -- so they are bit-identical across P-A and P-B except in the in-distribution regime, where P-A reads `train` out-of-fold and P-B reads the LODO 20% held-in slice (different eval subsets, not a LODO effect on the arm itself).

SPREAD GATE (sd >= 10 and bottom/top bin <= 0.80 on a 0-100 DV; 0-1 binary DVs rescaled), applied two ways. PARTIAL floor -- some datasets in a regime fail: those are DROPPED from the bar's mean and from the regime sub-label, and the surviving datasets carry the bar. Datasets dropped this way: evil/completely OOD: hprt, toxicchat, evil_pair. TOTAL floor -- EVERY dataset in a regime fails: the regime is still PLOTTED, from the floored data, but HATCHED + muted and EXCLUDED from the average panel. evil's generic chat is the only such regime (its sole dataset wildchat_rung is floored at sd 4.4 with 98.9% of mass in the bottom bin): its bars are shown so the regime is visibly measured-and-uninterpretable rather than silently absent, and they must NOT be read as effect sizes -- at that floor the rho is not estimable.

Sycophancy and hallucination lose nothing either way.

SEPARATELY -- SCOPE EXCLUSION, NOT A GATE FAILURE: evil/completely OOD: evil_tomgibbs. These datasets PASS the spread gate and were removed by an explicit user scope call, so evil's completely-OOD bar is now a SINGLE dataset (MHJ), not a mean over its OOD ladder. evil_tomgibbs is the excluded one and it is NOT neutral: ridge-on-mapped scored -0.399 (P-A) / -0.337 (P-B) there -- strongly NEGATIVE, and the only CI-disjoint context-vs-mapped comparison in the whole evil OOD set. Removing it therefore REMOVES the one setting where the mapped-answer readout was decisively worse than the context readout; read evil's OOD bar as 'MHJ only', never as evil OOD in general.

ERROR BARS -- one definition for every bar in both figures: +/-1 standard error of the PLOTTED MEAN, combining (i) within-cell sampling error over eval contexts, taken as half-width/1.96 of each cell's committed 95% bootstrap CI (the bootstrap resamples eval contexts), and (ii) between-cell spread. The combination rule follows whether the averaged cells share eval contexts. The three non-OOD regimes do: under P-B the K holdout fits all read ONE common pvsynth grid / WildChat slice / 20% held-in slice, so the context error is COMMON and does not average down -- var = v_within + s2_between/K. The completely-OOD regime does not (each cell is a different held-out corpus), nor do the three behaviours in the average panel, so there var = (v_within + tau2)/K with tau2 = max(0, s2_between - v_within) correcting the observed spread for the within-cell noise it already contains. A single-cell bar reduces to that cell's own bootstrap SE under both rules. This SUPERSEDES the previous mixed convention (bootstrap CI for single-dataset bars, bare s.e.m. across cells otherwise), whose s.e.m. branch read EXACTLY 0.000 for the persona-vector arms in every shared-context P-B regime -- those arms are not label-consuming, so the K LODO fits return literally one repeated number and their spread measures nothing; the bars looked precise where no precision had been estimated. Correspondingly, n_eval for a shared-context regime below now counts that ONE eval set, not K copies of it. Only the CONTEXT-based variant is scored (variant=context_end); the prefix-end arm is a stated scope deviation inherited from the parent round. Mapping is LINEAR throughout; no MLP or kernel arm.

METHOD ROSTER = the result2_fivemethod reference roster MINUS 'Ridge regression on real answer' (arm12_oracle_reg). That arm was NOT scored by the original P-A/P-B round -- that round ran five arms (arm1/arm4/arm6/arm7/arm11) and nothing ridge-on-the-real-answer -- and no re-score supplying it was found when this figure was built, so the arm is MISSING here rather than withheld. Nothing in this figure should be read as evidence about how ridge-on-the-real-answer performs.

result2_methods also carries arm8_map_ridge_true (ridge FIT on the real answer, APPLIED to the mapped answer), which appears in neither the reference figure nor this one.

n_eval per behaviour x regime --

evil: synthetic 200 [5/5 rungs kept]; generic chat 417 [5/5 rungs kept]; in-distribution 1,162 [5/5 rungs kept]; completely OOD 533 [1/4 rungs kept]

sycophancy: synthetic 200 [6/6 rungs kept]; generic chat 415 [6/6 rungs kept]; in-distribution 3,235 [6/6 rungs kept]; completely OOD 3,916 [6/6 rungs kept]

hallucination: synthetic 200 [2/2 rungs kept]; generic chat 411 [2/2 rungs kept]; in-distribution 3,350 [2/2 rungs kept]; completely OOD 7,188 [2/2 rungs kept]