


Judged behavior score in every evaluation setting: distribution and spread



Evil only: a judge REFUSAL verdict (the evaluated model declined the query) is scored 0 rather than dropped; sycophancy and hallucination are unchanged. Each row is labelled with the role it plays in the design, then the dataset it actually is. An in-distribution row is that behavior's FULL labeled pool, scored out-of-fold under 5 group-level folds — 'out-of-fold' is how the predictor is read, not a held-out data split. Hallucination's TriviaQA / NQ-Open / SimpleQA rows (above the dashed line) score fabrication rate x100, a different construct from the 0-100 trait rubric in every other row.