


Result 2 (older result2_methods protocol), five-method spread-flagged view: regression and persona-vector reads across evaluation regimes, hatched where the DV spread gate fails



Spearman rho of each predictor vs the judged behaviour-expression DV (graded 0-100 trait rubric; hallucination's in-distribution and completely-OOD cells instead score fabricated fraction x100 — a different construct, and the averaged panel mixes the two). PROTOCOL PROVENANCE: points come from the OLDER result2_methods protocol (result2_points.json), NOT the fair refit behind the committed four-panel figure — the fair round never ran the ridge-on-real-answer arm, so this is the only single-protocol source with all five methods; the companion old-vs-fair scatter shows how the two protocols compare on the four shared methods. The two 'real answer' methods are ORACLE reads (the source file labels them 'oracle'): they consume the model's ACTUAL generated answer at eval time; the context / mapped-answer methods score without it. All arms read CONTEXT-based representations only — prefix-end cells are excluded (recorded user-directed scope). Bar = mean of the committed per-replicate Spearman rho, frozen at the max-data operating slice (regime e1, full unlabeled pool, max labelled budget: evil 8k / sycophancy 16k / hallucination 16k); error bars = the committed bootstrap ci, frozen averaged over the same replicates (15 replicates on in-distribution + OOD cells; single replicate on synthetic + generic chat); an averaged bar carries the elementwise MEAN of its constituents' CI bounds — conservative, not a resampled CI of the mean. 'Completely OOD' = simple mean over the behaviour's OOD rungs: evil 2 (hh-rlhf red-team, ToxicChat), sycophancy 1 (held-out Reddit r/socialskills), hallucination 2 (NQ-Open, SimpleQA) — the average means different things per behaviour. Hatched muted groups FAIL the Result 1 spread gate (floor/ceiling mass <= 0.90 AND split-half reliability $r_{yy} \geq 0.50$ AND min detectable $\rho \leq 0.5 \times \sqrt{r_{yy}}$, per behaviour x setting): evil fails generic chat AND both OOD rungs, so those whole groups are untrustworthy (bars kept for completeness). Averaged panel EXCLUDES spread-failed cells — its generic-chat and OOD bars average 2 behaviours (sycophancy, hallucination), the rest 3. Per-bar n_{eval} — synthetic: 200 each; generic chat: 1,987/1,982/1,967 (evil/syco/hall); the FULL WildChat rung, unlike the fair figure's held-out ~20% split); in-distribution: 6,468 / 16,000 / 16,000; OOD rungs: evil 1,868 + 519, sycophancy 1,304, hallucination 3,167 + 4,021.