


Old-protocol vs fair-protocol Spearman rho, per (behavior, setting), four shared methods



Per-cell protocol-consistency check for the four methods scored under BOTH protocols: fair-refit Spearman rho (R2FAIR, result2_fair_points.json — source of the committed four-panel figure) vs the older result2_methods rho this five-method figure uses (ridge-on-real-answer exists only in the older file, so it cannot appear here). One point per (behavior, setting, method); 4 methods x 14 cells = 56 points; dashed line = y equals x. In the fair file the map arms split by map_kind (linear / mlp) while the older file does not — each older row is paired against the fair LINEAR row. Hollow markers = cells failing the Result 1 spread gate (all evil: WildChat, hh-rlhf red-team, ToxicChat). Known protocol differences besides the map/whitening refit: the older file's generic-chat cells score the FULL WildChat rung (n 1,967-1,987) while the fair file scores a held-out ~20% split (n 411-417).