




Input state: bare query -> answer (run only on the random-WildChat rung, so one setting wide)

Spearman rho, prediction vs judged behavior expression



reads the input state

- persona-vector projection @ bare query
- label-supervised direction @ bare query
- ridge @ bare query
- MLP @ bare query

reads through the fitted map

- persona-vector projection @ mapped answer
- ridge @ mapped answer (fit on predicted)
- ridge @ mapped answer (fit on real)
- map-pretrain then fine-tune
- stacked combiner

oracle (reads the real answer, not deployable)

- persona-vector projection @ real answer
- ridge @ real answer
- MLP @ real answer
- kernel ridge @ real answer

control / floor (not a usable method)

- identity + learned bias
- shuffled map
- shuffled pretraining
- text-embedding ridge
- surface features