




Input state: context (prefix + query) -> answer

Spearman rho, prediction vs judged behavior expression



evaluation setting, ordered by distance from the training distribution

colour: where the readout reads

- reads the input state directly
- reads through the fitted map to the answer state
- reads the real answer (oracle, not deployable)
- control / floor (not a usable method)

marker: readout type

- persona-vector projection
- label-supervised direction
- ridge
- ridge, fit on real answers
- MLP
- kernel ridge
- map-pretrain then fine-tune
- stacked combiner

controls (dotted, grey)

- control: identity + learned bias
- control: shuffled map
- control: shuffled pretraining
- control: text-embedding ridge
- control: surface features

reading the panel

- reliability ceiling (square root of r_{yy})
- spread gate failed (Result 1): ranking uninformative
- diagnostic column, not a deployment setting
- DV construct boundary (hallucination only)
- method not run beyond this setting