



Input state: prefix-end state -> answer

Spearman rho, prediction vs judged behavior expression


evil



sycophancy



hallucination



evaluation setting

reads the input state

- persona-vector projection @ prefix-end state
- label-supervised direction @ prefix-end state
- ridge @ prefix-end state
- MLP @ prefix-end state

reads through the fitted map

- persona-vector projection @ mapped answer
- ridge @ mapped answer (fit on predicted)
- ridge @ mapped answer (fit on real)
- map-pretrain then fine-tune
- stacked combiner

oracle (reads the real answer, not deployable)

- persona-vector projection @ real answer
- ridge @ real answer
- MLP @ real answer
- kernel ridge @ real answer

control / floor (not a usable method)

- identity + learned bias
- shuffled map
- shuffled pretraining
- text-embedding ridge
- surface features

Hallucination's held-out TriviaQA / NQ-Open / SimpleQA settings (left of the dashed line) score fabrication rate x100, a different construct from the 0-100 trait rubric everywhere else.