




Result 2: persona-vector projection onto context -> mapped answer -> real answer



reads the context

■ context

reads through the fitted context -> answer map

■ mapped answer (linear map)

■ mapped answer (kernel map)

■ mapped answer (MLP map)

reads the real answer (ceiling)

■ real answer

reading the marks

⋯ dashed edge: map fit on the GENERIC WildChat pool only (no ADD cell at that setting)

○ open circle: single replicate (CI is a within-draw bootstrap over eval contexts)

Missing bar = that method was not run at that setting (the two nonlinear maps ran at the out-of-fold train rung only); never a zero. Hallucination's in-distribution / NQ-Open / SimpleQA settings (right of the dashed line) score fabrication rate x100, a different construct from the 0-100 trait rubric everywhere else.