




LAYOUT: within each evaluation-regime group, bars are clustered by INPUT SOURCE — context / mapped answer / real answer (short labels context / mapped / real beneath each cluster; whitespace separates the clusters); within a cluster the LIGHTER bar is the fitted ridge regression and the DARKER bar the persona-vector projection. Spearman rho of each predictor vs the judged behaviour-expression DV (graded 0-100 trait rubric; hallucination's in-distribution and completely-OOD cells instead score fabricated fraction x100 — a different construct, and the averaged panel mixes the two). PROTOCOL PROVENANCE: points come from the OLDER result2_methods protocol (result2_points.json), NOT the fair refit behind the committed four-panel figure — the fair round never ran the ridge-on-real-answer arm, so this is the only single-protocol source with all five methods. The two 'real answer' methods are ORACLE reads (the source file labels them 'oracle'): they consume the model's ACTUAL generated answer at eval time; the context / mapped-answer methods score without it. All arms read CONTEXT-based representations only — prefix-end cells are excluded (recorded user-directed scope). Bar = mean of the committed per-replicate Spearman rho_frozen at the max-data operating slice (regime e1, full unlabeled pool, max labelled budget: evil 8k / sycophancy 16k / hallucination 16k); error bars = the committed bootstrap ci_frozen averaged over the same replicates (15 replicates on in-distribution + OOD cells; single replicate on synthetic + generic chat); an averaged bar carries the elementwise MEAN of its constituents' CI bounds — conservative, not a resampled CI of the mean. 'Completely OOD' = simple mean over the behaviour's OOD rungs: evil 2 (hh-rlhf red-team, ToxicChat), sycophancy 1 (held-out Reddit r/socialskills), hallucination 2 (NQ-Open, SimpleQA) — the average means different things per behaviour. MUTED (faded) groups FAIL the Result 1 spread gate (floor/ceiling mass ≤ 0.90 AND split-half reliability $r_{yy} \geq 0.50$ AND min detectable rho $\leq 0.5 \times \sqrt{r_{yy}}$, per behaviour x setting): evil fails generic chat AND both OOD rungs, so those whole groups are untrustworthy (bars kept for completeness). Averaged panel EXCLUDES spread-failed cells — its generic-chat and OOD bars average 2 behaviours (sycophancy, hallucination), the rest 3. The generic-chat column UNDERSTATES the regression arms: the older grid scored the FULL WildChat rung (n 1,987/1,982/1,967 evil/syco/hall) while the fair refit scores a held-out ~20% split (n 411-417). Per-bar n_eval elsewhere — synthetic: 200 each; in-distribution: 6,468 / 16,000 / 16,000; OOD rungs: evil 1,868 + 519, sycophancy 1,304, hallucination 3,167 + 4,021. SIX-BAR VARIANT: the mapped-answer cluster carries BOTH mapped-answer regression arms — one ridge fitted on the mapped answer and applied to the mapped answer (the committed figure's 'Ridge regression on mapped answer'), and one ridge fitted on the REAL answer and applied to the mapped answer (oracle answers at fit time, mapped answers at eval time). Both have full 14-cell coverage in the source file.