


Result 5c: the two composition channels crossed — eliciting data in the MAP pool vs in the LABEL pool



f_U is the fraction of the UNLABELED map pool that is trait-eliciting rather than generic WildChat; f_L is the same fraction for the LABELED readout pool. Both are composition at FIXED size -- eliciting rows REPLACE generic rows -- so neither axis is confounded with quantity. Cells come from the main grid's compose family at an unlabeled pool of 5,000, evaluated on the in-distribution rung. Realized coverage: 8 (f_U, f_L, L) cells for evil ONLY -- the compose family was never run for sycophancy, hallucination, the ($f_U=0, f_L=1$) corner was never run at all, and ($f_U=0.5, f_L=0$) has no $L=8,000$ cell, so a missing point is missing data rather than a null. Points are the mean over the cell's replicates, error bars $1.96 \times \text{SEM}$.