


Result 5a: does a bigger UNLABELED generic map pool buy behavior prediction?



Each panel is one behavior x evaluation rung. x is the number of unlabeled generic WildChat context->answer pairs the MAP was fit on; the labeled readout budget is held at each behavior's maximum (evil 8,000; sycophancy and hallucination 16,000). Points are the mean frozen-layer Spearman rho over the source's 5 label draws x 3 seeds, error bars 1.96 x SEM over those 15 replicates. Arms with no row at a given U are omitted rather than interpolated. Rungs are the ORIGINAL roster this grid was run on -- the newer MHJ / PAIR / tom-gibbs / SycophancyEval rungs postdate it and appear in the LODO figure instead. Slice: context-based variant, synthetic (E1) extraction regime.