


Result 5b: replacing judged ELICITING labels with judged GENERIC labels, budget held fixed



Composition at FIXED budget, never addition: the labeled readout pool is 1,500 rows throughout, and x is the fraction of those rows drawn from judged random WildChat rather than the behavior's eliciting corpus, so a change along x cannot be confounded with quantity. Points are the mean frozen-layer Spearman rho over the source's replicates at that cell; error bars 1.96 x SEM (a single-replicate cell is drawn with a zero-width bar, which means no replicate spread was estimable, not that the estimate is exact). The label-free projection arms (persona vector on context / mapped / real answer) do not consume labels, so their curves are expected to be flat in x -- they are drawn as the reference the label-consuming arms move against.