




Result 5d: leave-one-dataset-out — what does the readout lose on the dataset it never saw?



Under the leave-one-dataset-out protocol the readout trains on an 80% group-level slice of every eliciting dataset EXCEPT one, which is held out whole; there is one fit per held-out dataset (evil 5, sycophancy 6, hallucination 2). Each bar pools that fit's evaluation cells into three classes: the dataset it never saw, the 20% held-in slices of the sibling eliciting datasets it did train on, and random WildChat. The synthetic persona-vector grid belongs to none of the three and is excluded. ONLY the label-consuming arms are drawn -- the persona-vector projections do not train a readout, so which dataset left the fit cannot move them, and showing them would imply an effect the design cannot produce. Bars are the mean over the pooled cells with a 1.96 x SEM half-width; per-bar cell counts are in the JSON sidecar.