



Never-trained prefixes amplify map change with length, but stay below the trained conversation prefix

A — never-trained rungs (circles) vs trained-prefix anchors

raw trained-vs-base map difference m
(median per-context norm, layer 19)



B — all 36 cells: D per rung, 95% CIs; | = bare



● never-trained rung (round 4) ● swapped trained prefix (round 3)
 ★ own trained prefix (round 3) ● bare prompts

— sycophancy — persona-trained — casual writing — persona-trained
 — impoliteness — persona-trained — sycophancy — conversation-trained

● short rung (10 tok)
 ● mid rung (85 tok)
 ● long rung (730 tok)
 | bare (round 3, n-matched)