

Every arm plotted individually (16 arms per behavior: 4 trained contexts × 2 regimes × 2 seeds); shaded region = never-trained personas

