

Predicting the per-context write $w(x)$ from the base context vector $c_0(x)$ — layer 19. Variance-weighted R^2 vs the test mean of w , so $R^2 > 0$ is context-dependent predictability beyond the static mean write.

