

Every estimator on every metric: brighter = better within each column

context arm (963k train, layer 19, pool 1000)

constant train-mean	-0.04	0.80	0.00	0.00	0.01	500
identity copy	-2.54	0.64	0.25	0.45	0.34	9.00
identity + bias	-0.92	0.73	0.53	0.65	0.62	1.00
linear map (ridge)	0.75	0.95	0.81	0.85	0.85	1.00
neural map (w=8192)	0.81	0.96	0.84	0.87	0.88	1.00
neural map (w=32768)	0.81	0.97	0.85	0.88	0.89	1.00
kernel map (Nystrom)	0.81	0.96	0.84	0.87	0.88	1.00
	pooled R2	mean cosine	acc@1 (euclid)	acc@1 (CSLS)	MRR	median rank

prefix arm (LOFO, n~43 train, layer 18, pool 50)

constant fold-mean	-0.33	0.98	0.00	0.00	0.05	30
identity copy	-45.10	0.71	0.10	0.20	0.26	6.00
scaled identity	-17.74	0.71	0.04	0.20	0.14	15
per-dim rescale	-2.33	0.95	0.30	0.30	0.43	4.00
identity + bias	-6.53	0.92	0.84	0.74	0.91	1.00
LOFO ridge	0.54	0.99	0.04	0.16	0.25	6.50
LOFO neural map	-7.84	0.88	0.04	0.04	0.22	7.00
	pooled R2	mean cosine	acc@1 (euclid)	acc@1 (CSLS)	MRR	median rank