

Every estimator on every metric — context arm, Qwen2.5-7B-Instruct layer 19, candidate pool 1,000 held-out LMSYS rows. Brighter = better WITHIN each column.
◀ marks a baseline already in use; row-label color = training rows (blue 963,444 / orange 3,600 / purple 50).

linear map (ridge) [963k]	0.754	0.617	0.955	0.274	0.805	0.800	0.848	0.850	1.000
neural map (w=8192) [963k]	0.810	0.688	0.965	0.295	0.841	0.833	0.873	0.885	1.000
neural map (w=32768) [963k]	0.813	0.699	0.965	0.294	0.850	0.845	0.883	0.892	1.000
kernel map (Nystrom RBF) [963k]	0.807	0.698	0.964	0.289	0.838	0.835	0.872	0.880	1.000
identity + learned bias [963k]	-0.920	-2.881	0.727	0.155	0.532	0.567	0.652	0.617	1.000
identity copy [963k] ◀	-2.539	-4.084	0.637	0.126	0.254	0.250	0.446	0.340	9.000
constant train-mean [963k]	-0.045	-0.016	0.798	0.000	0.001	0.001	0.001	0.007	500
linear map (ridge) [3600]	0.705	0.551	0.945	0.261	0.720	0.720	0.802	0.790	1.000
identity + learned bias [3600]	-0.865	-2.830	0.726	0.158	0.503	0.553	0.638	0.600	1.000
per-dim rescale (diagonal) [3600] ◀	0.096	0.026	0.827	0.098	0.088	0.116	0.352	0.153	52
scaled identity (one scalar) [3600] ◀	-0.694	-0.519	0.637	0.126	0.079	0.250	0.446	0.139	62
constant train-mean [3600]	-0.001	-0.001	0.807	-0.000	0.001	0.001	0.001	0.007	500
linear map (ridge) [n=50]	0.384	0.212	0.883	0.146	0.065	0.066	0.142	0.142	30
identity + learned bias [n=50]	-0.890	-2.873	0.725	0.157	0.501	0.537	0.639	0.598	1.000
	pooled R^2	per-dim R^2 (median)	mean cosine (raw)	mean cosine – shuffled null	acc@1 (euclidean)	acc@1 (cosine)	acc@1 (CSLS)	MRR (euclidean)	median rank (euclidean)