


Is the assistant a privileged persona? Assistant chat map re-used on each story character


#2054 lattice · story answer boundary = bare label · AI-likeness = judge-scored (claude-sonnet-4-5, k=5, n≈300) on each character's OWN on-policy answers · user turn excluded · y-axis shared within a row only

held-out R^2 (context arm, layer 19)

Qwen2.5-7B-Instruct · inserted answer text
n=4,135–4,945 paired rows per character




Qwen2.5-7B-Instruct · on-policy answer text
n=3,481–3,647 paired rows per character




held-out R^2 (context arm, layer 19)

Qwen2.5-7B (base) · inserted answer text
n=4,135–4,945 paired rows per character



Qwen2.5-7B (base) · on-policy answer text
n=3,481–3,647 paired rows per character



— control: → the ASSISTANT in story, persona held (inserted target)
— control: character → character, median of 12 pairs

→ Helios (AI-likeness 72) → Wren (AI-likeness 62)
→ Vex (AI-likeness 55) → Dana (AI-likeness 53)