


Re-using the assistant's chat-template context→answer map under a changed framing


#2054 lattice · read = boundary before the assistant starts speaking → mean over answer tokens · every panel $n_{train} > d=3,584$ (ambient basis)

held-out R^2 (context arm, layer 19)


Qwen2.5-7B-Instruct · INSERTED (controlled arm)
n=11,901 paired rows · cross-render n=11,901 (4-way shared)




Qwen2.5-7B base · INSERTED (controlled arm)
n=11,901 paired rows · cross-render n=11,901 (4-way shared)



Qwen2.5-7B-Instruct · ON-POLICY (joint arm)
n=8,000 paired rows · cross-render n=8,000 (4-way shared)



Qwen2.5-7B base · ON-POLICY (joint arm)
n=7,999 paired rows · cross-render n=7,999 (4-way shared)



transfer tier (more correction allowed →)

transfer tier (more correction allowed →)

- chat → bare text
- ▲ chat → in story (bare-label boundary)
- · cross-render fit: chat context → target answer (not a tier)
- chat → in story (attributed-quote boundary)
- target map's own within-cell R^2 (ceiling)
- shuffle null (95th pct)

CONTROLLED arm (top row): the answer text is held fixed across framings, so a delta is attributable to framing.

JOINT arm (bottom row): the answer is regenerated in each framing, so a delta mixes what is said with how it is encoded — it is NOT a framing effect.

The dash-dot cross-render line is NOT a tier: every tier re-uses the FROZEN chat operator, while this refits a fresh $d \times d$ ridge on (chat context, target answer) pairs.

With full d^2 capacity it upper-bounds how much the chat context DETERMINES the target answer — it cannot show the two framings share an operator.