


Re-using the assistant's chat-template context→answer map under a changed framing


#2054 lattice · read = boundary before the assistant starts speaking → mean over answer tokens · every panel $n_{train} > d=3,584$ (ambient basis)

held-out R^2 (context arm, layer 19)

Qwen2.5-7B-Instruct · INSERTED (controlled arm)
n=11,901 paired rows




Qwen2.5-7B base · INSERTED (controlled arm)
n=11,901 paired rows




held-out R^2 (context arm, layer 19)

Qwen2.5-7B-Instruct · ON-POLICY (joint arm)
n=8,000 paired rows



Qwen2.5-7B base · ON-POLICY (joint arm)
n=7,999 paired rows



transfer tier (more correction allowed →)

transfer tier (more correction allowed →)

● chat → bare text

▲ chat → in story (bare-label boundary)

■ shuffle null (95th pct)

■ chat → in story (attributed-quote boundary)

.... target map's own within-cell R^2 (ceiling)

CONTROLLED arm (top row): the answer text is held fixed across framings, so a delta is attributable to framing.

JOINT arm (bottom row): the answer is regenerated in each framing, so a delta mixes what is said with how it is encoded — it is NOT a framing effect.