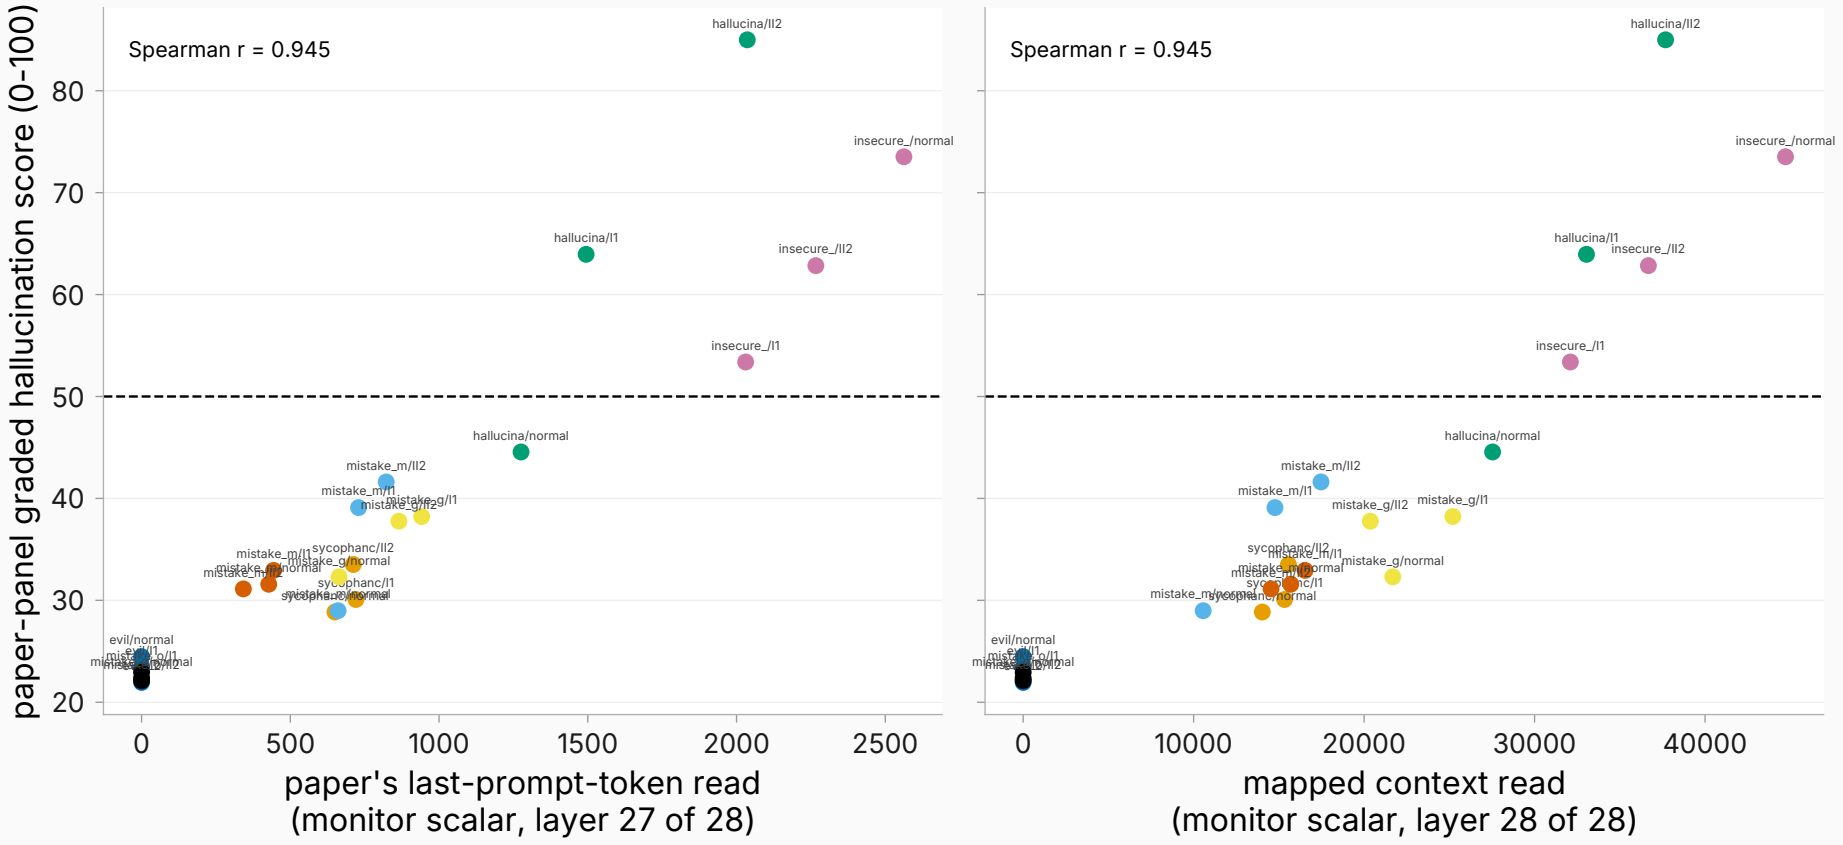
Spearman r = 0.945
paper-panel graded hallucination score (0-100)
paper's last-prompt-token read
(monitor scalar, layer 27 of 28)
Spearman r = 0.945
mapped context read
(monitor scalar, layer 28 of 28)
hallucina/ll2
insecure_/normal
hallucina/l1
insecure_/ll2
insecure_/l1
hallucina/normal
mistake_m/ll2
mistake_m/l1
mistake_g/ll2
mistake_g/l1
sycophanc/ll2
mistake_g/normal
sycophanc/l1
evil/normal