


4a dataset-level screening vs post-ft trait shift — hallucination

post-ft trait score (graded)


raw

$\rho = -0.2696$ (jk -0.4247--0.0091, n=24/8 fam)




exact_dp

$\rho = -0.4065$ (jk -0.474--0.1506, n=24/8 fam)




prompttoken_dp

$\rho = -0.3678$ (jk -0.5364--0.074, n=24/8 fam)




mapped_dp_context

$\rho = -0.2583$ (jk -0.4273--0.0857, n=24/8 fam)




mapped_dp_prefix

$\rho = -0.2696$ (jk -0.4247--0.0091, n=24/8 fam)




probe_diff_context

$\rho = 0.1948$ (jk 0.1948--0.6325, n=24/8 fam)



probe_diff_prefix

$\rho = 0.3383$ (jk 0.2039--0.3987, n=24/8 fam)



family

- evil
- hallucination
- insecure_code
- mistake_gsm8k
- mistake_math
- mistake_medical
- mistake_opinions
- sycophancy