


4a dataset-level screening vs post-ft trait shift — sycophancy

post-ft trait score (graded)

raw


$\rho=0.0304$ (jk -0.1532–0.226, n=24/8 fam)



dataset-level screening score

exact_dp


$\rho=0.0661$ (jk -0.0688–0.2714, n=24/8 fam)



dataset-level screening score

prompttoken_dp


$\rho=0.0852$ (jk -0.3416–0.0675, n=24/8 fam)



dataset-level screening score

mapped_dp_context


$\rho=0.0304$ (jk 0.3039–0.5286, n=24/8 fam)



dataset-level screening score

mapped_dp_prefix


$\rho=0.0304$ (jk -0.1532–0.226, n=24/8 fam)



dataset-level screening score

probe_diff_context


$\rho=0.0026$ (jk 0.0403–0.3974, n=24/8 fam)



dataset-level screening score

probe_diff_prefix

$\rho=0.0087$ (jk 0.1688–0.6104, n=24/8 fam)



dataset-level screening score

family

- evil
- hallucination
- insecure_code
- mistake_gsm8k
- mistake_math
- mistake_medical
- mistake_opinions
- sycophancy