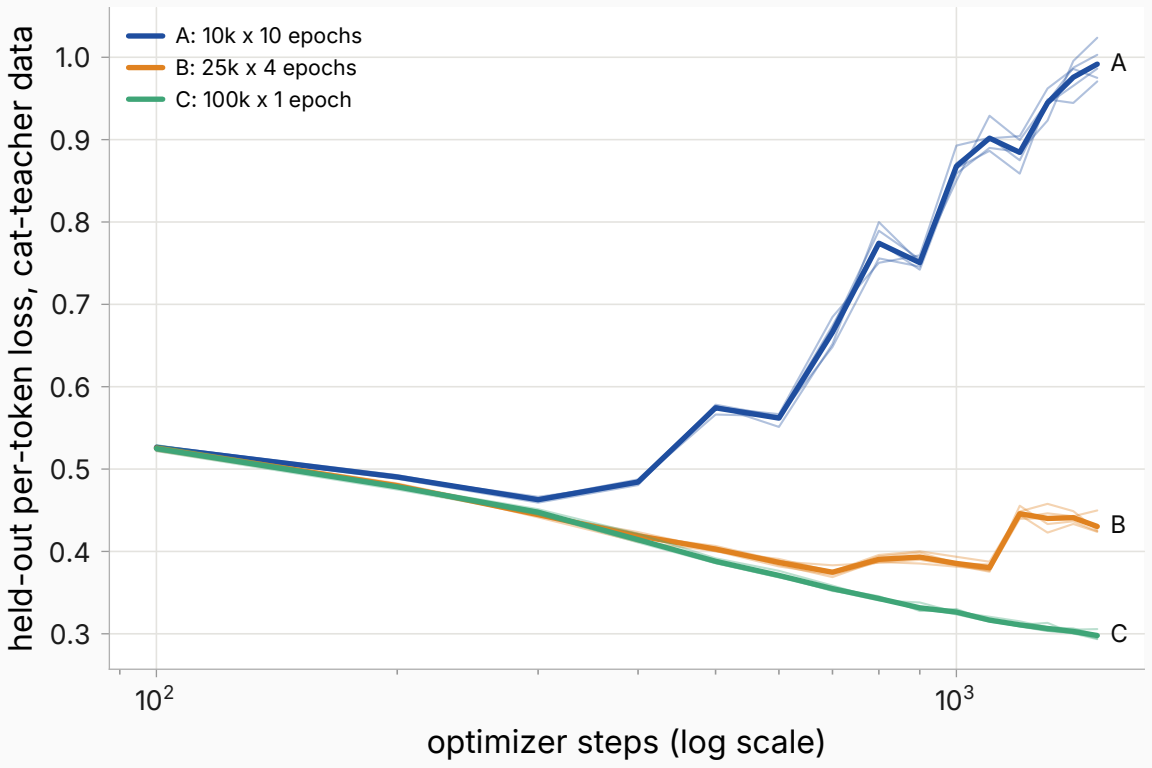
A: 10k x 10 epochs
B: 25k x 4 epochs
C: 100k x 1 epoch
held-out per-token loss, cat-teacher data
1.0
0.9
0.8
0.7
0.6
0.5
0.4
0.3
A
B
C
10^2
10^3
optimizer steps (log scale)