


FIT UNIT

- Pooled (four corpora)
- MATH
- ▲ GSM8K train
- ▼ ContextHub
- ◆ MMLU
- Post-SFT model's own map (retention 1)

QWEN2.5-7B-INSTRUCT → OPENTHINKER3-7B, LAYER 19, POOLED AND PER-CORPUS FITS

Only a change of basis on both sides recovers the map after reasoning SFT

Retention (transferred R^2 / own R^2) ↑



Correction fit on training folds before applying the pre-SFT map to the post-SFT model