

BASELINE

— Boundary token $\rightarrow$ next sentence

▲ Encoder (BGE)

▼ Copy + bias


PREDICTOR

—●— Linear

—◆— Nonlinear

A · 50,000 TRAINING CONTEXTS

Predictability across layers



QWEN2.5-7B-INSTRUCT

METRIC

—●— Held-out R^2

—○— Top-1 retrieval

B · LAYER 19, 10,000-CANDIDATE RETRIEVAL

Scaling with training data

